


You have to sweat the details!

[Michael Hay](#)



A detailed view of Yokohama's skyline, photo taken by Michael Hay circa 2023

I vividly remember saying the words in this title. Sitting in a room with management consultants and corporate leadership, these words came out of my mouth full of feeling. I was compelled to speak but hesitant and nervous all at once. The context and company aren't important to this narrative, but the words and the feeling behind them are. I say this because being a believer in the career path, general mission and daily plight of Products teams details matter and sweating them is important. Too often very senior Products leaders ask for conclusions quickly resulting in shoddy and indefensible work that doesn't bring the team along. It leaves you wondering how can this be helped?

Introducing Caffeine

While [John Goodman](#) and I are huge coffee drinkers this Caffeine is not what

you find in a cup of joe. It is instead our Machine Intelligence Inferencing Engine built to help Products teams do their best work. The core capabilities of Caffeine cover the following functionalities.

- Abstractive summarization of large documents using Generative AI.
- Automated tag discovery for documents and Company corpuses using Automated Term Extraction.
- Per document topic modeling with BERT and topic labeling with Generative AI.
- Document and Company corpus similarity analysis using Sentence Transformers.

When Caffeine is enabled for Mediumroast for GitHub, it affords users the following features, immediately available, to help Products teams speed up discovery work.

- Reveal guided “so what” points for each document in your collection. These points include keywords, relevant excerpts, and a sentence that ideally looks like a requirement statement.
- For every Company, within the Repository, understand which other Company is most similar and which of its documents are most/least interesting. This can be helpful, based upon content you’ve captured, to become the basis for competitive analysis.

Below is a screenshot of report present in Mediumroast’s GitHub discovery repository because we’re bootstrapping our competitive analysis with our own tool. Here [Aha!](#) is found to be the most competitively similar to [Atlassian](#) for the set of companies we have in the repository. We’ve used our tooling to help us pinpoint, of the analyzed competitive field, which company or companies warrant the most review; John and I will relate this in a later article. Until then the utility of Mediumroast for GitHub, with Caffeine enabled, is to help us more efficiently sweat the right details.

Doing business

Our idealism and inspiration from [Rational Software](#)'s business model leads us to do business with our tools, methods and services to [battle disinformation](#). Therefore, we're starting business using our tooling in a services wrapper to help anyone interested in accelerating competitive discovery work. We're already seeing utility and want to sharpen both our tool and our methodology with early customers. If you're interested contact us at hello@mediumroast.io.

The lagniappe, some crow?

In a previous article, "[Startups should be in the public cloud?](#)", I wrote on

using an on-prem system for our Machine Learning and Generative AI work. Well that system powers Caffeine and we moved out of the garage and into the closet. Further John and I have been selectively moving small non-AI components to Microsoft's Azure. I'm fairly certain, given how assertive public cloud enthusiasts were, an enthusiast will try and send John and I crow to eat. If you're interested here's a YouTube video of [someone cooking crow](#), and it looks yummy. So bring it on!

